



Approximate  
Nearest Neighbor  
Search via Group  
Testing

Authors

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

# APPROXIMATE NEAREST NEIGHBOR SEARCH VIA GROUP TESTING

**Saksham Rathi**  
(22B1003)

**Kshitij Vaidya**  
(22B1829)

**Ekansh Ravi Shankar**  
(22B1032)

CS754: ADVANCED IMAGE PROCESSING  
UNDER PROF. AJIT RAJWADE

Indian Institute of Technology Bombay  
Spring 2025

# Contents



Approximate  
Nearest Neighbor  
Search via Group  
Testing

Authors

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- 1 Introduction
- 2 Locality Sensitive Hashing
- 3 Distance-Sensitive Bloom Filters
- 4 Algorithm
- 5 Theory

# Nearest Neighbor Search



Approximate  
Nearest Neighbor  
Search via Group  
Testing

Authors

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- Nearest neighbor search is a fundamental problem with many applications in machine learning systems.
- **Task:** Given a dataset  $D = \{x_1, x_2, \dots, x_N\}$ , the goal is to build a data structure that can be queried with any point  $q$  to obtain a small set of points  $x_i \in D$  that have high similarity (low distance) to the query. This structure is called an index.
- Such tasks frequently arise in genomics, web-scale data mining, machine learning, and other large-scale applications.

# Group Testing



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- We are given a set  $D$  of  $N$  items, with  $k$  positives (“hits”) and  $N - k$  negatives (“misses”).
- **Goal:** Identify all positive items using fewer than  $N$  group tests.
- A *group test* is positive iff at least one item in the group is positive.
- **Testing Variants:** Can be *noisy* (with false positives/negatives), *adaptive* (tests depend on previous results), or *non-adaptive* (all tests run in parallel).
- The paper uses a **doubly regular design**: Each item appears in an equal number of tests; each test has an equal number of items.

# Formal Problem Statement



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- **( $R, c$ )-Approximate Near Neighbor:** Given a dataset  $D$ , if there exists a point within distance  $R$  of a query  $y$ , return some point within distance  $c \cdot R$ , with high probability.
  - $R$  is the distance threshold (radius).
  - $c > 1$  is the approximation factor.
- Any algorithm that solves the randomized nearest neighbor problem also solves the approximate near neighbor problem with  $c = 1$  and any  $R \geq$  distance to the nearest neighbor.
- (Definition) **Randomized Nearest neighbor:** Given a dataset  $D$  and a distance metric  $d(\cdot, \cdot)$  and a failure probability  $\delta \in [0, 1]$ , construct a data structure which, given a query point  $y$  reports the point  $x \in D$  with the smallest distance  $d(x, y)$  with probability greater than  $1 - \delta$ .

# Locality Sensitive Hashing



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

A hash function  $h(x) \rightarrow \{1, \dots, R\}$  is a function that maps an input  $x$  to an integer in the range  $[1, R]$ .

The two points  $x$  and  $y$  are said to collide if  $h(x) = h(y)$ .

$$s(x, y) = \Pr_H(h(x) = h(y))$$

For now, we will assume that  $s(x, y) = \text{sim}(x, y)$ .

For any positive integer  $L$ , we may transform an LSH family  $H$  with collision probability  $s(x, y)$  into a new family having  $s(x, y)^L$  by sampling  $L$  hash functions from  $H$  and concatenating the values to obtain a new hash code  $[h_1(x), h_2(x), \dots, h_L(x)]$ . If the original hash family had the range  $[1, R]$ , the new hash family has the range  $[1, R^L]$ .

# Locality Sensitive Hashing



Approximate  
Nearest Neighbor  
Search via Group  
Testing

Authors

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- **Locality Sensitive Hashing (LSH)** algorithms use an LSH function to partition the dataset into buckets.
- The hash function is selected so that the distance between points in the same bucket is likely to be small.
- To find the near neighbors of a query, we hash the query and compute the distance to every point in the corresponding bucket.
- **Count-Based LSH** identifies neighbors by simply counting how many times two points land in the same hash bucket across multiple hash functions.

# Distance-Sensitive Bloom Filters



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- (Definition) **Approximate Set Membership:** Given a set  $D$  of  $N$  points and similarity thresholds  $S_L$  and  $S_H$ , construct a data structure which, given a query point  $y$ , has:  
True Positive Rate: If there is  $x \in D$  with  $\text{sim}(x, y) > S_H$ , the structure returns true w.p.  $\geq p$   
False Negative Rate: If there is no  $x \in D$  with  $\text{sim}(x, y) > S_L$ , the structure returns true w.p.  $\leq q$
- The distance-sensitive Bloom filter solves this problem using LSH functions and a 2D bit array. The structure consists of  $m$  binary arrays that are each indexed by an LSH function. There are three parameters: the number of arrays  $m$ , a positive threshold  $t \leq m$ , and the number of concatenated hash functions  $L$  used within each array.



# Distance-Sensitive Bloom Filters



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

- To construct the filter, we insert elements  $x \in D$  by setting the bit located at array index  $[m, h_m(x)]$  to 1.
- To query the filter, we determine the  $m$  hash values of the query  $y$ . If at least  $t$  of the corresponding bits are set, we return true. Otherwise, we return false.
- (Theorem) Assuming the existence of an LSH family with collision probability  $s(x, y) = \text{sim}(x, y)$ , the distance-sensitive Bloom filter solves the approximate membership query problem with

$$p \geq 1 - \exp(-2m(-t + S_H^L)^2)$$

$$q \leq \exp(-2m(-t + NS_L^L)^2)$$



**Input:** Dataset  $D$  of size  $N$ , positive integers  $B$  and  $R$ , similarity threshold  $S$

**Output:** A FLINNG search index consisting of membership sets  $M_{r,b}$  and group tests  $C_{r,b}$

- For  $r = 0$  to  $R - 1$ :
  - Let  $\pi(D)$  be a random permutation of  $D$
  - For  $b = 0$  to  $B - 1$ :
    - Define  $M_{r,b} = \{\pi(D)_i \mid i \bmod B = b\}$
- For  $r = 0$  to  $R - 1$ :
  - For  $b = 0$  to  $B - 1$ :
    - Construct a classifier  $C_{r,b}$  for membership set  $M_{r,b}$  with true positive rate  $p$  and false negative rate  $q$



- If we apply a similarity threshold to the dataset, we obtain a near neighbor set  $K = \{x \in D | \text{sim}(x, y) \geq S\}$ . We consider  $K$  to be the set of “positives” in the group testing problem.
- In order to do so, we split the dataset  $D$  into a set of groups, which we visualize as a  $B \times R$  grid of cells. Each cell has a group of items  $M_{r,b}$  and a corresponding group test  $C_{r,b}$ . To assign items to cells, we evenly distribute the  $N$  points among the  $B$  cells in each column of the grid, and we independently repeat this assignment process  $R$  times.



**Input:** A FLINNG index and a query  $y$

**Output:** Approximate set  $\hat{K}$  of neighbors with similarity greater than the threshold  $S$

- Initialize  $\hat{K} = \{1, \dots, N\}$
- For  $r = 0$  to  $R - 1$ :
  - Initialize  $Y = \emptyset$
  - For  $b = 0$  to  $B - 1$ :
    - If  $C_{r,b}(y) = 1$  then:  $Y = Y \cup M_{r,b}$
  - $\hat{K} = \hat{K} \cap Y$



- To query the index with a point  $y$ , we begin by querying each classifier. If  $C_{r,b}(y) = 1$ , then at least one of the points in  $M_{r,b}$  has high similarity to  $y$ . We collect all of these “candidate points” by taking the union of the  $M_{r,b}$  sets for which  $C_{r,b}(y) = 1$ .
- We repeat this process for each of the  $R$  repetitions to obtain  $R$  candidate sets, one for each column in the grid.
- With high probability, each candidate set contains the true neighbors, but it may also have some non-neighbors that were included in  $M_{r,b}$  by chance. To filter out these points, we intersect the candidate sets to obtain our approximate near neighbor set  $\hat{K}$ .

# Group Testing: Runtime and Accuracy



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

[Introduction](#)

[Locality Sensitive  
Hashing](#)

[Distance-  
Sensitive Bloom  
Filters](#)

[Algorithm](#)

[Theory](#)

**Lemma:** Suppose we have a dataset  $D$  of points, where a subset  $K \subseteq D$  is “positive” and the rest are “negative”. Construct a  $B \times R$  grid of tests, where each test has i.i.d true positive rate  $p$  and false negative rate  $q$ . Then the algorithm reports points as “positive” with probability:

$$\Pr(\text{Report } x \mid x \in K) \geq p^R$$

$$\Pr(\text{Report } x \mid x \notin K) \leq \left[ q \left( \frac{eN(B-1)}{B(N-1)} \right)^{|K|} + p \left( 1 - \left( \frac{N(B-1)}{eB(N-1)} \right)^{|K|} \right) \right]^R$$

# Group Testing: Runtime and Accuracy



Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory

The cost of group testing inference includes the cost to do all  $B \times R$  tests, plus the cost of intersecting the positive groups.

**Theorem:** Under the assumptions in the previous lemma, let us suppose that each test runs in  $\mathcal{O}(T)$ . Then with probability  $1 - \delta$ :

$$t_{\text{query}} = \mathcal{O} \left( BRT + \frac{RN}{B} (p|K| + qB) \log\left(\frac{1}{\delta}\right) \log N \right)$$

# Bounding the Test Cost



To distinguish between the  $|K|$  nearest neighbors and the rest of the dataset:

$$S_H = \text{sim}(x_{|K|}, y) = s_{|K|}$$

$$S_L = \text{sim}(x_{|K|+1}, y) = s_{|K|+1}$$

(Definition)  **$\gamma$ -Stable query:** We say that the query is  $\gamma$ -stable if

$$\frac{\log(s_{|K|})}{\log(s_{|K|}) - \log(s_{|K|+1})} \leq \gamma$$

**Theorem:** Given a true positive rate  $p$ , false negative rate  $q$  and stability parameter  $\gamma$ , it is possible to choose  $m$ ,  $L$  and  $t$  so that the resulting distance-sensitive Bloom filter has true positive rate  $p$  and false negative rate  $q$  for all  $\gamma$ -stable queries. The query time is

$$\mathcal{O}(mL) = \mathcal{O}(-\log(\min(q, 1 - p))N\gamma\log(N))$$

Approximate  
Nearest Neighbor  
Search via Group  
Testing

[Authors](#)

Introduction

Locality Sensitive  
Hashing

Distance-  
Sensitive Bloom  
Filters

Algorithm

Theory